

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [5]: df = pd.read_csv(r"C:\Users\Roshan Ramdas Kate\Downloads\Social_Network_Ads.csv")
```

```
In [6]: df.shape
```

```
Out[6]: (400, 5)
```

```
In [7]: df.head()
```

```
Out[7]:   User ID  Gender  Age  EstimatedSalary  Purchased
0    15624510    Male    19        19000          0
1    15810944    Male    35        20000          0
2    15668575  Female    26        43000          0
3    15603246  Female    27        57000          0
4    15804002    Male    19        76000          0
```

```
In [8]: df.dtypes
```

```
Out[8]: User ID      int64
Gender        object
Age         int64
EstimatedSalary  int64
Purchased     int64
dtype: object
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: User ID      0
Gender        0
Age         0
EstimatedSalary  0
Purchased     0
dtype: int64
```


```
In [10]: df.describe()
```

Out[10]:

| | User ID | Age | EstimatedSalary | Purchased |
|--------------|--------------|------------|-----------------|------------|
| count | 4.000000e+02 | 400.000000 | 400.000000 | 400.000000 |
| mean | 1.569154e+07 | 37.655000 | 69742.500000 | 0.357500 |
| std | 7.165832e+04 | 10.482877 | 34096.960282 | 0.479864 |
| min | 1.556669e+07 | 18.000000 | 15000.000000 | 0.000000 |
| 25% | 1.562676e+07 | 29.750000 | 43000.000000 | 0.000000 |
| 50% | 1.569434e+07 | 37.000000 | 70000.000000 | 0.000000 |
| 75% | 1.575036e+07 | 46.000000 | 88000.000000 | 1.000000 |
| max | 1.581524e+07 | 60.000000 | 150000.000000 | 1.000000 |


```
In [11]: g = sns.catplot(x = "Gender",y = "Purchased",data = df,kind = "bar",height = 4)
g.set_ylabels("Purchased Probability")
plt.show
```

Out[11]: <function matplotlib.pyplot.show(close=None, block=None)>




```
In [12]: M2 = pd.crosstab(df.Gender, df.Purchased, normalize='index')
print(M2)
M2.plot.bar(figsize=(6,4),stacked=True)
plt.legend(title='Gender vs Purchased', loc='upper right')
plt.show()
```

| Purchased | 0 | 1 |
|-----------|----------|----------|
| Gender | | |
| Female | 0.622549 | 0.377451 |
| Male | 0.663265 | 0.336735 |



```
In [14]: corr = df.corr()
print(corr.shape)
plt.figure(figsize=(8,8))
sns.heatmap(corr, cbar=True, square= True, fmt=' .1f', annot=True, annot_kws={'size':(4, 4)})
Out[14]: <AxesSubplot:>
```



```
In [15]: X=df.drop(['Gender','Purchased'],axis=1)
Y= df['Purchased']
X.head()
```

Out[15]:

| | User ID | Age | EstimatedSalary |
|---|----------|-----|-----------------|
| 0 | 15624510 | 19 | 19000 |
| 1 | 15810944 | 35 | 20000 |
| 2 | 15668575 | 26 | 43000 |
| 3 | 15603246 | 27 | 57000 |
| 4 | 15804002 | 19 | 76000 |

```
In [17]: # Split the dataset into training and testing datasets
from sklearn.model_selection import train_test_split
# Shuffle and split the data into training and testing subsets
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_st
```

```
In [19]: from sklearn.linear_model import LogisticRegression
# instantiate the model (using the default parameters)
basemodel= LogisticRegression()
# fit the model with data
```

```
basemodel.fit(X_train,y_train)
print("Training accuracy:", basemodel.score(X_train,y_train)*100)
```

Training accuracy: 78.75

```
In [20]: y_predict= basemodel.predict(X_test)
print("Testing accuracy:", basemodel.score(X_test,y_test)*100)
```

Testing accuracy: 73.75

```
In [21]: from sklearn.metrics import accuracy_score
Acc=accuracy_score(y_test,y_predict)
print(Acc)
```

0.7375

```
In [22]: from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test,y_predict)
print(cm)
```

```
[[49  3]
 [18 10]]
```

```
In [23]: from sklearn.metrics import precision_recall_fscore_support
prf= precision_recall_fscore_support(y_test,y_predict)
print('precision:',prf[0])
print('Recall:',prf[1])
print('fscore:',prf[2])
print('support:',prf[3])
```

precision: [0.73134328 0.76923077]

Recall: [0.94230769 0.35714286]

fscore: [0.82352941 0.48780488]

support: [52 28]

```
In [24]: from sklearn.metrics import classification_report
cr= classification_report(y_test,y_predict)
print(cr)
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.73 | 0.94 | 0.82 | 52 |
| 1 | 0.77 | 0.36 | 0.49 | 28 |
| accuracy | | | 0.74 | 80 |
| macro avg | 0.75 | 0.65 | 0.66 | 80 |
| weighted avg | 0.74 | 0.74 | 0.71 | 80 |

```
In [ ]:
```